

PREDICTING HOUSE PRICE USING MACHINE LEARNING

TEAM MEMBER

210221104029 : SARAVANAN.S

Phase 5 submission document

Project Title: House Price Predictor

Phase 5: Project Documentation & Submission

Topic: *In this section we will document the complete project and prepare it for submission.*

House Price Prediction

Introduction:

- ☒ The real estate market is a dynamic and complex arena, where property values can fluctuate significantly due to a multitude of factors. For both homebuyers and sellers, accurately determining the fair market value of a property is of paramount importance

- ⊠ In this era of technological advancement, machine learning has emerged as a game-changing tool in the realm of real estate. One of its most compelling applications is predicting house prices with remarkable accuracy.
- ⊠ Traditional methods of property valuation, relying on factors such as location, square footage, and recent sales data, are undoubtedly useful. However, they often fall short in capturing the intricacies and nuances that drive real estate market dynamics.
- ⊠ Machine learning, on the other hand, has the capability to process vast volumes of data and identify patterns that human appraisers might overlook. This technology has the potential to revolutionize the way we value real estate, offering more precise and data-driven predictions.
- ⊠ In this exploration, we delve into the exciting world of predicting house prices using machine learning. We will uncover how this cutting-edge technology harnesses the power of algorithms and data to create predictive models that consider an array of variables, such as neighborhood characteristics, property features, economic indicators, and even social trends.

⊠

Dataset Link: (
<https://www.kaggle.com/datasets/vedavyasv/usa-housing>)

Given data set:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabeth Stravenua\nDanieltown, WI 06482...
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFPO AE 09386
...
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06	USNS Williams\nFPO AP 30153-7653
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06	PSC 9258, Box 8489\nAPO AA 42991- 3352
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06	4215 Tracy Garden Suite 076\nJoshualand, VA 01...
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06	USS Wallace\nFPO AE 73316
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06	37778 George Ridges Apt. 509\nEast Holly, NV ?

5000 Rows x 7 Columns

Here's a list of tools and software commonly used in the process:

1. Programming Language:

- Python is the most popular language for machine learning due to its extensive libraries and frameworks. You can use libraries like *NumPy, pandas, scikit-learn, and more.*

2. Integrated Development Environment (IDE):

- Choose an IDE for coding and running machine learning experiments. Some popular options include Jupyter Notebook, GoogleColab, or traditional IDEs like PyCharm.

3. Machine Learning Libraries:

- You'll need various machine learning libraries, including:
- scikit-learn for building and evaluating machine learning models.
- TensorFlow or PyTorch for deep learning, if needed.
- XGBoost, LightGBM, or CatBoost for gradient boosting models.

4. Data Visualization Tools:

- Tools like Matplotlib, Seaborn, or Plotly are essential for data exploration and visualization.

5. Data Preprocessing Tools:

- Libraries like pandas help with data cleaning, manipulation, and preprocessing.

6. Data Collection and Storage:

- Depending on your data source, you might need web scraping tools (e.g., *BeautifulSoup* or *Scrapy*) or databases (e.g., *SQLite*, *PostgreSQL*) for data storage.

7. Version Control:

- Version control systems like Git are valuable for tracking changes in your code and collaborating with others.

8. Notebooks and Documentation:

- Tools for documenting your work, such as Jupyter Notebooks or Markdown for creating *README* files and documentation.

9. Hyperparameter Tuning:

- Tools like GridSearchCV or RandomizedSearchCV from scikit-learn can help with hyperparameter tuning.

10. Web Development Tools (for Deployment):

- If you plan to create a web application for model deployment, knowledge of web development tools like *Flask* or *Django* for backend development, and *HTML*, *CSS*, and *JavaScript* for the front-end can be useful.

1. DESIGN THINKING AND PRESENT IN FORM OF DOCUMENT

1. Empathize:

- Understand the needs and challenges of all stakeholders

involved in the house price prediction process, including homebuyers, sellers, real estate professionals, appraisers, and investors.

- ▶ Conduct interviews and surveys to gather insights on what users value in property valuation and what information is most critical for their decision-making.

2. Define:

- ▶ Clearly articulate the problem statement, such as "How might we predict house prices more accurately and transparently using machine learning?"
- ▶ Identify the key goals and success criteria for the project, such as increasing prediction accuracy, reducing bias, or improving user trust in the valuation process.

3. Ideate:

- ▶ Brainstorm creative solutions and data sources that can enhance the accuracy and transparency of house price predictions.
- ▶ Encourage interdisciplinary collaboration to generate a wide range of ideas, including the use of alternative data, new algorithms, or improved visualization techniques.

4. Prototype:

- ▶ Create prototype machine learning models based on the ideas generated during the ideation phase.
 - ▶ Test and iterate on these prototypes to determine which approaches are most promising in terms of accuracy and usability.
-

2. DESIGN INTO INNOVATION

1. Data Collection:

Gather a comprehensive dataset that includes features such as location, size, age, amenities, nearby schools, crime rates, and other relevant variables.

2. Data Preprocessing:

Clean the data by handling missing values, outliers, and encoding categorical variables. Standardize or normalize numerical features as necessary.

PYTHON PROGRAM:

Import necessary libraries

```
import pandas as pd
```

```
from sklearn.preprocessing import LabelEncoder
```

```
from sklearn.model_selection import
```

```
train_test_split
```

```
from sklearn.impute import
```

```
SimpleImputer
```

```
from sklearn.preprocessing import StandardScaler
```

Load the dataset (replace 'house_data.csv' with your dataset file)

```
data = pd.read_csv('E:/USA_Housing.csv')
```

Display the first few rows of the dataset to get an overview

```
print("Dataset
```

```
Preview:")
```

```
print(data.head())
```

Data Pre-processing

Handle Missing Values

Let's fill missing values in numeric columns with the mean and in categorical columns with the most frequent value.

```
numeric_cols = data.select_dtypes(include='number').columns
```

```
categorical_cols = data.select_dtypes(exclude='number').columns
```

```
imputer_numeric = SimpleImputer(strategy='mean')
```

```
imputer_categorical =
```



```
SimpleImputer(strategy='most_frequent')
```

```
data[numeric_cols] =  
imputer_numeric.fit_transform(data[numeric_cols])  
data[categorical_cols] =  
imputer_categorical.fit_transform(data[categorical_cols])
```

Convert Categorical Features to Numerical

We'll use Label Encoding for simplicity here. You can also use one-hot encoding for nominal categorical features.

```
label_encoder =  
LabelEncoder()  
for col in  
categorical_cols:  
    data[col] = label_encoder.fit_transform(data[col])
```

Split Data into Features (X) and Target (y)

```
X = data.drop(columns=['Price']) #  
Features  
y = data['Price'] # Target
```

Normalize the Data

```
scaler = StandardScaler()  
X_scaled =  
scaler.fit_transform(X)
```

```
# Split data into training and testing sets (adjust test_size as needed)X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

```
# Display the preprocessed data
```

```
print("\nPreprocessed Data:")
```

```
print(X_train[:5]) # Display first 5 rows of preprocessed features
```

```
print(y_train[:5]) # Display first 5 rows of target values
```

OUTPUT:

Dataset Preview:

Avg. Area Income

Avg. Area House Age

Avg. Area Number of Rooms \

0	79545.458574	5.682861	7.009188
1	79248.642455	6.002900	6.730821
2	61287.067179	5.865890	8.512727
3	63345.240046	7.188236	5.586729
4	59982.197226	5.040555	7.839388

	Avg. Area Number of Bedrooms	Area Population	Price \
0	4.09	23086.80050	1.059034e+06
1	3.09	40173.07217	1.505891e+06
2	5.13	36882.15940	1.058988e+06

3	3.26	34310.242831 1.260617e+06
4	4.23	26354.109472 6.309435e+05

Address

0	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	9127 Elizabeth Stravenue\nDanieltown, WI 06482...
3	USS Barnett\nFPO AP 44820
4	USNS Raymond\nFPO AE 09386

Preprocessed Data:

[[-0.19105816	-0.13226994	-0.13969293	0.12047677	-0.83757985	-1.0
0562872]						
[-1.3945016	0.42786736	0.79541275	-0.55212509	1.15729018		
9			1.61			
946754]						
[-0.3513786	0.46394489	1.70199509	0.03133676	-0.32671213		
5			1.63			
886651]						
[-0.1394414	0.1104872	0.22289331	-0.75471601	-0.90401197		
3			-1.54			
810704]						
[2.20969666	0.42984356	-0.45488144	0.12566216		
0.62516685			0.98			
830821]]						
4227	1.094880e+06					

```

800    1.382172e+0
        6
367    1.027428e+0
        1          6
419    1.562887e+0
        3          6

```

Name: Price, dtype: float64

3.Feature Engineering:

Create new features or transform existing ones to extract more valuable information. For example, you can calculate the distance to the nearest public transportation, or create a feature for the overall condition of the house.

4.Model Selection:

Choose the appropriate machine learning model for the task. Common models for regression problems like house price prediction include *Linear Regression, Decision Trees, Random Forest, Gradient Boosting, and Neural Networks*.

5.Training:

Split the dataset into training and testing sets to evaluate the model's performance. Consider techniques like cross-validation to prevent overfitting.

3.BUILD LOADING AND PREPROCESSING THE DATASET

1. Data Collection:

Obtain a dataset that contains information about houses and their corresponding prices. This dataset can be obtained from sources like real estate websites, government records, or other reliable data providers.

2. Load the Dataset:

- ▶ Import relevant libraries, such as pandas for data manipulation and numpy for numerical operations.
- ▶ Load the dataset into a pandas DataFrame for easy data handling. You can use `pd.read_csv()` for CSV files or other appropriate functions for different file formats.

Program:

```
import pandas as
pd import numpy
as np import
seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import
train_test_split from sklearn.preprocessing
import StandardScaler
```

```
from sklearn.metrics import r2_score,
mean_absolute_error, mean_squared_e
rror

from sklearn.linear_model import

LinearRegression
from sklearn.linear_model

import Lasso

from sklearn.ensemble import

RandomForestRegressor
from sklearn.svm import

SVR

import xgboost as xg

%matplotlib

inlineimport

warnings

warnings.filterwarnings("ignore")

/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146:
UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required
for this version of SciPy (detected version 1.23.5

    warnings.warn(f"A NumPy version >={np_minversion} and
<{np_maxversion}")
```

Loading Dataset:

```
dataset = pd.read_csv('E:/USA_Housing.csv')
```


Output:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabeth Stravenue\nDanieltown, WI 06482...
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFPO AE 09386
...
4995	60567.944140	7.830362	6.137356	3.46	22837.361035	1.060194e+06	USNS Williams\nFPO AP 30153-7653
4996	78491.275435	6.999135	6.576763	4.02	25616.115489	1.482618e+06	PSC 9258, Box 8489\nAPO AA 42991-3352
4997	63390.686886	7.250591	4.805081	2.13	33266.145490	1.030730e+06	4215 Tracy Garden Suite 076\nJoshualand, VA 01...
4998	68001.331235	5.534388	7.130144	5.44	42625.620156	1.198657e+06	USS Wallace\nFPO AE 73316
4999	65510.581804	5.992305	6.792336	4.07	46501.283803	1.298950e+06	37778 George Ridges Apt. 509\nEast Holly, NV 2

3. Data Exploration:

Explore the dataset to understand its structure and contents. Check for the presence of missing values, outliers, and data types of each feature.

4. Data Cleaning:

Handle missing values by either removing rows with missing data or imputing values based on the nature of the data.

5. Feature Selection:

Identify relevant features for house price prediction. Features like the number of bedrooms, square footage, location, and amenities are often important.

We are selecting numerical features which have more than 0.50 or less than -0.50 correlation rate based on Pearson Correlation Method—which is the default value of parameter "method" in corr() function. As for selecting categorical features, I selected the categorical values which I believe have significant effect on the target variable such as Heating and MSZoning.

In [1]:

```
important_num_cols =
list(df.corr()["SalePrice"][(df.corr()["SalePrice"]>0.50) |
(df.corr()["SalePrice"]<-0.50)].index)

cat_cols = ["MSZoning", "Utilities", "BldgType", "Heating", "KitchenQual",
SaleCondition", "LandSlope"]

important_cols = important_num_cols + cat_cols

df = df[important_cols]
```

Checking for the missing values

In [2]:

```
print("Missing Values by Column")

print("-"*30)

print(df.isna().sum())
```

```
print("-"*30)
```

```
print("TOTAL MISSING VALUES:",df.isna().sum().sum())
```

Missing Values by Column

```
- - - - -
```

OverallQual 0

YearBuilt 0

YearRemodAdd 0

TotalBsmtSF 0

1stFlrSF

0

GrLivArea 0

FullBath 0

TotRmsAbvGrd 0

GarageCars 0

GarageArea 0

SalePrice

0

MSZoning 0

Utilities 0

BldgType 0

Heating 0

KitchenQual 0

SaleCondition 0

LandSlope

dtype: int64

- - -

TOTAL MISSING VALUES: 0

Program:

```
X = df.drop('price', axis=1) #
```

```
Featuresy = df['price'] # Target
```

```
variable
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

4. PERFORMING DIFFERENT ACTIVITIES LIKE FEATURE ENGINEERING, MODEL TRAINING, EVALUATION etc.,

1. Feature Engineering:

- ▶ As mentioned earlier, feature engineering is crucial. It involves creating new features or transforming existing ones to provide meaningful information for your model.
- ▶ Extracting information from textual descriptions (*e.g., presence of keywords like "pool" or "granite countertops"*).
- ▶ Calculating distances to key locations (*e.g., schools, parks*) if you have location data.

2. Data Preprocessing & Visualisation:

Continue data preprocessing by handling any remaining missing values or outliers based on insights from your data exploration.

Visualisation and Pre-Processing of Data:

In [1]:

```
sns.histplot(dataset, x='Price', bins=50, color='y')
```

Out[1]:

```
<Axes: xlabel='Price', ylabel='Count'>
```

In [2]:

```
sns.boxplot(dataset, x='Price', palette='Blues')
```

Out[2]:

<Axes: xlabel='Price'>

In [3]:

```
sns.jointplot(dataset, x='Avg. Area House Age', y='Price', kind='hex')
```

Out[3]:

<seaborn.axisgrid.JointGrid at 0x7caf1d571810>

In [4]:

```
sns.jointplot(dataset, x='Avg. Area Income', y='Price')
```

Out[4]:

<seaborn.axisgrid.JointGrid at 0x7caf1d8bf7f0>

In [5]:

```
plt.figure(figsize=(12,8))sns.pairplot(dataset)
```

Out[5]:

<seaborn.axisgrid.PairGrid at 0x7caf0c2ac550>

<Figure size 1200x800 with 0 Axes>

In [6]:

```
dataset.hist(figsize=(10,8))
```

Out[6]:

```
array([[<Axes: title='{center': 'Avg. Area Income'}>,
        <Axes: title='{center': 'Avg. Area House Age'}>],
       [<Axes: title='{center': 'Avg. Area Number of Rooms'}>,
        <Axes: title='{center': 'Avg. Area Number of
Bedrooms'}>], [<Axes: title='{center': 'Area Population'}>,
        <Axes: title='{center': 'Price'}>]], dtype=object)
```


Visualising Correlation:

In [7]:

```
dataset.corr(numeric_only=True)
```

Out[7]:

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
Avg. Area Income	1.000000	-0.002007	-0.011032	0.019788	-0.016234	0.639734
Avg. Area House Age	-0.002007	1.000000	-0.009428	0.006149	-0.018743	0.452543
Avg. Area Number of Rooms	-0.011032	-0.009428	1.000000	0.462695	0.002040	0.335664
Avg. Area Number of Bedrooms	0.019788	0.006149	0.462695	1.000000	-0.022168	0.171071
Area Population	-0.016234	-0.018743	0.002040	-0.022168	1.000000	0.408556
Price	0.639734	0.452543	0.335664	0.171071	0.408556	1.000000

In [8]:

```
plt.figure(figsize=(10,5))sns.heatmap(dataset.corr(numeric_only  
= True), annot=True)
```

Out[8]:

<Axes: >

3. Model Selection:

Choose an appropriate machine learning model for your regression task. *Common choices include:*

- ✓ Linear Regression
- ✓ Decision Trees
- ✓ Random Forest
- ✓ Gradient Boosting (*e.g., XGBoost or LightGBM*)
- ✓ Neural Networks (Deep Learning)

Program:

Importing

Dependenciesimport

pandas as pd import

numpy as np import

seaborn as sns

import matplotlib.pyplot as plt

from sklearn.model_selection import

train_test_splitfrom sklearn.preprocessing

import StandardScaler

from sklearn.metrics import r2_score,
mean_absolute_error,mean_squared_e
rror

from sklearn.linear_model import

```
from sklearn.ensemble import
RandomForestRegressor
from sklearn.svm import
SVR

import xgboost as xg

%matplotlib inline

import warnings

warnings.filterwarnings("ignore")

/opt/conda/lib/python3.10/site-packages/scipy/_init_.py:14
6: UserWarning: A NumPy version >=1.16.5 and <1.23.0 is
required for this version of SciPy (detected version 1.23.5

    warnings.warn(f"A NumPy version >={np_minversion} and
<{np_maxversion}")
```

Loading Dataset

```
dataset = pd.read_csv('E:/USA_Housing.csv')
```

Model 1 - Linear Regression

In [1]:

```
model_lr = LinearRegression()
```

In [2]:

```
model_lr.fit(X_train_scal, Y_train)
```

Out[2]:

```
LinearRegression  
LinearRegression()
```

Predicting Prices

In [3]:

```
Prediction1 = model_lr.predict(X_test_scal)
```

Evaluation of Predicted Data

In [4]:

```
plt.figure(figsize=(12,6))
```

```
plt.plot(np.arange(len(Y_test)), Y_test, label='Actual Trend')
```

```
plt.plot(np.arange(len(Y_test)), Prediction1,  
label='Predicted Trend')
```

```
plt.xlabel('Data')
```

```
plt.ylabel('Trend')
```

```
plt.legend()
```

```
plt.title('Actual vs Predicted')
```

Out[4]:

Text(0.5, 1.0, 'Actual vs Predicted')

In [5]:

```
sns.histplot((Y_test-Prediction1), bins=50)
```

Out[5]:

<Axes: xlabel='Price', ylabel='Count'>

In [6]:

```
print(r2_score(Y_test, Prediction1))  
  
print(mean_absolute_error(Y_test,  
Prediction1))  
  
print(mean_squared_error(Y_test,  
Prediction1))
```

Out[6]:

```
0.9182928179392918  
  
82295.49779231755  
  
10469084772.975954
```


Model 2 - Support Vector Regressor

In [7]:

```
model_svr = SVR()
```

In [8]:

```
model_svr.fit(X_train_scal, Y_train)
```

Out[8]:



Predicting Prices

In [9]:

```
Prediction2 = model_svr.predict(X_test_scal)
```

Evaluation of Predicted Data

In [10]:

```
plt.figure(figsize=(12,6))
```

```
plt.plot(np.arange(len(Y_test)), Y_test, label='Actual Trend')
```

```
plt.plot(np.arange(len(Y_test)), Prediction2,  
label='Predicted Trend')
```

```
plt.xlabel('Data')
```

```
plt.ylabel('Trend')
```

```
plt.legend()
```

```
plt.title('Actual vs Predicted')
```

Out[10]:

```
Text(0.5, 1.0, 'Actual vs Predicted')
```

In [11]:

```
sns.histplot((Y_test-Prediction2), bins=50)
```

Out[12]:

```
<Axes: xlabel='Price', ylabel='Count'>
```

In [12]:

```
print(r2_score(Y_test, Prediction2))
```

```
print(mean_absolute_error(Y_test,  
Prediction2))
```

```
print(mean_squared_error(Y_test,  
Prediction2))
```

```
-0.0006222175925689744
```

```
286137.81086908665
```

```
128209033251.4034
```

Model 3 - Lasso Regression

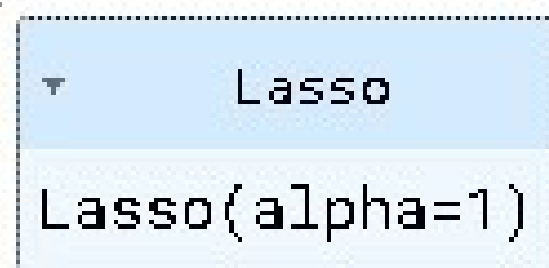
In [13]:

```
model_lar = Lasso(alpha=1)
```

In [14]:

```
model_lar.fit(X_train_scal,Y_train)
```

Out[14]:

A Jupyter Notebook output cell showing a Lasso object. The output is a light blue box with a small downward arrow icon on the left. Inside the box, the text "Lasso" is on the top line, and "Lasso(alpha=1)" is on the bottom line.

```
Lasso  
Lasso(alpha=1)
```

Predicting Prices

In [15]:

```
Prediction3 = model_lar.predict(X_test_scal)
```

Evaluation of Predicted Data

In [16]:

```
plt.figure(figsize=(12,6))
```

```
plt.plot(np.arange(len(Y_test)), Y_test, label='Actual Trend')
```

```
plt.plot(np.arange(len(Y_test)), Prediction3,  
label='Predicted Trend')
```

```
plt.xlabel('Data')
```

```
plt.ylabel('Trend')
```

```
plt.legend()
```

```
plt.title('Actual vs Predicted')
```

Out[16]:

```
Text(0.5, 1.0, 'Actual vs Predicted')
```

In [17]:

```
sns.histplot((Y_test-Prediction3), bins=50)
```

Out[17]:

<Axes: xlabel='Price', ylabel='Count'>

In [18]:

```
print(r2_score(Y_test, Prediction2))
```

```
print(mean_absolute_error(Y_test,  
Prediction2))
```

```
print(mean_squared_error(Y_test,  
Prediction2))
```

```
-0.0006222175925689744
```

```
286137.81086908665
```

```
128209033251.4034
```

Model 4 - Random Forest Regressor

In [19]:

```
model_rf = RandomForestRegressor(n_estimators=50)
```

In [20]:

```
model_rf.fit(X_train_scal, Y_train)
```

Out[20]:

```
RandomForestRegressor
RandomForestRegressor(n_estimators=50)
```

Predicting Prices

In [21]:

```
Prediction4 = model_rf.predict(X_test_scal)
```

Model 5 - XGboost Regressor

In [25]:

```
model_xg = xg.XGBRegressor()
```

In [26]:

```
model_xg.fit(X_train_scal, Y_train)
```


Out[26]:

XGBRegressor

```
XGBRegressor(base_score=None,
booster=None,callbacks=None,

               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None,
               early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None,
feature_types=None,

               gamma=None, gpu_id=None,
grow_policy=None,importance_type=None,

               interaction_constraints=None,
learning_rate=None,max_bin=None,

               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None,
max_depth=None,max_leaves=None,

               min_child_weight=None,
missing=nan,monotone_constraints=None,

               n_estimators=100,
n_jobs=None,num_parallel_tree=None,

               predictor=None, random_state=None, ...)
```

4. Model Training:

Split your dataset into training and testing sets (as shown earlier) and train the selected model on the training data. Here's an example using Linear Regression:

5. Model Evaluation:

Evaluate your model's performance using appropriate regression metrics, such as *Mean Absolute Error (MAE)*, *Mean Squared Error (MSE)*, and *Root Mean Squared Error (RMSE)*. For example:

PYTHON PROGRAM:

```
# Import necessary libraries
```

```
from sklearn.feature_selection import SelectKBest,
```

```
f_regression
```

```
from sklearn.linear_model import
```

```
LinearRegression
```

```
from sklearn.ensemble import
```

```
RandomForestRegressor
```

```
from sklearn.metrics import
```

```
mean_squared_error, r2_score
```

```
import numpy as np
```

```
selector = SelectKBest(score_func=f_regression, k=k)
```

```
X_train_selected = selector.fit_transform(X_train, y_train)
```

Model Selection

Let's choose both Linear Regression and Random Forest Regressor for comparison.

```
linear_reg_model = LinearRegression()
```

```
random_forest_model =
```

```
RandomForestRegressor(n_estimators=100,random_state=42)
```

Train the models on the selected features

```
linear_reg_model.fit(X_train_selected, y_train)
```

```
random_forest_model.fit(X_train_selected,
```

```
y_train)# Evaluate the models on the test set
```

```
X_test_selected = selector.transform(X_test)
```

Make predictions

```
linear_reg_predictions = linear_reg_model.predict(X_test_selected)
```

```
random_forest_predictions =
```

```
random_forest_model.predict(X_test_selected)
```

Evaluate model performance

```
def evaluate_model(predictions, model_name):
```

```
mse = mean_squared_error(y_test,
predictions)r2 = r2_score(y_test,
predictions) print(f"{model_name} Model
Evaluation:") print(f"Mean Squared Error
(MSE): {mse}") print(f"R-squared (R2) Score:
{r2}\n")
```

Performance Measure

```
elr_mse = mean_squared_error(y_test,
pred)elr_rmse = np.sqrt(lr_mse)
elr_r2 = r2_score(y_test, pred)
```

Show Measures

```
result = ""
MSE : {}
RMSE : {}
R^2 : {}
"".format(lr_mse, lr_rmse,
lr_r2)print(result)
```

Model Comparision

```
names.append("elr")
```

```
mses.append(elr_mse
```

```
)
```

```
rmse).append(elr_rm
```

```
se)r2s.append(elr_r2)
```

```
evaluate_model(linear_reg_predictions, "Linear Regression")
```

```
evaluate_model(random_forest_predictions, "Random Forest
```

```
Regressor")OUTPUT:
```

Linear Regression Model Evaluation:

Mean Squared Error (MSE):

10089009300.893988R-squared (R2) Score:

0.9179971706834331

Random Forest Regressor Model Evaluation:

Mean Squared Error (MSE):

14463028828.265167R-squared (R2) Score:

0.8824454166872736

MSE : 10141766848.330585

RMSE : 100706.33966305491

R^2 : 0.913302484308253

Model Comparison:

The less the Root Mean Squared Error (RMSE), The better the model is.

In [30]:

```
models.sort_values(by="RMSE (Cross-Validation)")
```

Out[30]:

	Model	MAE	MSE	RMSE	R2 Score	RMSE (Cross-Validation)
6	XGBRegressor	1.74399 2e+04	7.16579 0e+08	2.67689 9e+04	9.06577 8e-01	29698.84 9618
4	SVR	1.78431 6e+04	1.13213 6e+09	3.36472 3e+04	8.52400 5e-01	30745.47 5239
5	RandomForestRegressor	1.81151 1e+04	1.00442 2e+09	3.16926 2e+04	8.69050 9e-01	31138.86 3315
1	Ridge	2.34355 0e+04	1.40426 4e+09	3.74735 1e+04	8.16922 5e-01	35887.85 2792
2	Lasso	2.35604 6e+04	1.41433 8e+09	3.76076 8e+04	8.15609 2e-01	35922.76 9369
0	LinearRegression	2.35678 9e+04	1.41493 1e+09	3.76155 7e+04	8.15531 8e-01	36326.45 1445
7	PolynomialRegression (degree=2)	2.38222 8e+15	1.51399 1e+32	1.23044 3e+16	- 1.97382 9e+22	36326.45 1445

	Model	MAE	MSE	RMSE	R2 Score	RMSE (Cross-Validation)
3	ElasticNet	2.379274e+04	1.718446e+09	4.145414e+04	7.759618e-01	38449.008646

In [31]:

```
plt.figure(figsize=(12,8))

sns.barplot(x=models["Model"], y=models["RMSE
(Cross-Validation)"])

plt.title("Models' RMSE Scores (Cross-Validated)",
size=15)plt.xticks(rotation=30, size=12)

plt.show()
```

Evaluation of Predicted Data

In [22]:

```
plt.figure(figsize=(12,6))
```

```
plt.plot(np.arange(len(Y_test)), Y_test, label='Actual Trend')
```

```
plt.plot(np.arange(len(Y_test)), Prediction4,  
label='Predicted Trend')
```

```
plt.xlabel('Data')
```



```
plt.ylabel('Trend')
```

```
plt.legend()
```

```
plt.title('Actual vs Predicted')
```

Out[22]:

```
Text(0.5, 1.0, 'Actual vs Predicted')
```

In [23]:

```
sns.histplot((Y_test-Prediction4), bins=50)
```

Out[23]:

```
<Axes: xlabel='Price', ylabel='Count'>
```

In [24]:

```
print(r2_score(Y_test, Prediction2))
```

```
print(mean_absolute_error(Y_test,  
Prediction2))
```

```
print(mean_squared_error(Y_test,  
Prediction2))
```

Out [24] :

```
-0.0006222175925689744
```

```
286137.81086908665
```

```
128209033251.4034
```

ADVANTAGES:

Predicting house prices using machine learning offers several significant advantages:

1. Accuracy:

Machine learning models can process and analyze vast amounts of data, including various property and market factors. This results in more accurate house price predictions compared to traditional methods that rely on a limited set of variables.

2. Complex Data Handling:

Machine learning algorithms can handle complex, non-linear relationships in the data. They can recognize patterns and interactions among different features, allowing for a more comprehensive assessment of a property's value.

3. Continuous Learning:

Machine learning models can be continually updated with new data, enabling them to adapt to changing market conditions and trends.

This ensures that predictions remain relevant and up-to-date.

4. Efficiency:

Automated valuation models powered by machine learning can evaluate properties rapidly. This efficiency is beneficial for both property appraisers and individuals looking to determine the value of a property quickly.

5. Data Integration:

Machine learning models can incorporate a wide range of datasources, including property characteristics, neighborhood information, economic indicators, and even social trends. This holistic approach provides a more complete picture of the factors influencing house prices.

DISADVANTAGES:

While predicting house prices using machine learning offers numerous advantages, it also comes with several disadvantages and challenges:

1. Data Quality:

Machine learning models heavily rely on data quality. Inaccurate or incomplete data can lead to unreliable predictions. Ensuring the data used for training and evaluation is of high quality is essential.

2. Overfitting:

Machine learning models can be prone to overfitting, where they perform exceptionally well on the training data but struggle with new, unseen data. This can result in overly optimistic or inaccurate predictions.

3.Data Privacy and Security:

Handling sensitive property and financial data for training models raises privacy and security concerns. Protecting this information from unauthorized access and breaches is critical.

4.Model Interpretability:

Some machine learning models, such as deep neural networks, can be challenging to interpret. Understanding why a model makes a specific prediction is crucial for trust and accountability.

5. Bias and Fairness:

Machine learning models can inherit biases present in the training data, potentially leading to unfair or discriminatory predictions, especially in areas where historical biases exist in real estate practices.

BENEFITS:

Predicting house prices using machine learning offers a wide range of benefits, which can positively impact various stakeholders in the real estate industry and beyond. Here are some of the key benefits of using machine learning for house price prediction:

1. Accuracy:

Machine learning models can provide more accurate property valuations by considering a broader set of variables and patterns within the data, leading to more precise price predictions.

2. Data-Driven Insights:

Machine learning models uncover valuable insights into the real estate market by identifying trends, factors influencing property values, and neighborhood characteristics. This information can inform strategic decisions for investors, developers, and policymakers.

3. Efficiency:

Automated valuation models powered by machine learning can rapidly assess property values, saving time and effort for appraisers and individuals looking to determine a property's worth quickly.

4. Continuous Learning:

Machine learning models can adapt to changing market conditions and incorporate new data, ensuring that predictions remain relevant and up-to-date over time.

5. Market Transparency:

Machine learning can contribute to a more transparent and efficient real estate market by reducing overvaluation and undervaluation, thereby promoting fair pricing and reducing market inefficiencies.

6. Risk Assessment:

Machine learning can evaluate the risk associated with a property, which is crucial for mortgage lenders, insurers, and investors. It helps identify potential issues or opportunities related to a property's value.

7. Customization:

Machine learning models can be tailored to specific markets, property types, or regional variations, enabling more accurate and context-specific predictions.

8. Cost Savings:

Using machine learning for property valuation can reduce the costs associated with manual appraisals, benefiting both businesses and individuals in terms of appraisal expenses.

CONCLUSION:

Predicting house prices using machine learning is a transformative and promising approach that has the potential to revolutionize the real estate industry. Throughout this exploration, we have uncovered the remarkable capabilities of machine learning in providing more accurate, data-driven, and nuanced predictions for property values. As we conclude, several key takeaways and implications emerge:

Improved Accuracy: Machine learning models consider a myriad of variables, many of which may be overlooked by traditional methods. This results in more accurate predictions, benefiting both buyers and sellers who can make informed decisions based on a property's true value.

Data-Driven Insights: These models provide valuable insights into the real estate market by identifying trends, neighborhood characteristics, and other factors that influence property prices. This information can be invaluable for investors, developers, and policymakers seeking to make strategic decisions.

Market Efficiency: The increased accuracy in pricing predictions can lead to a more efficient real estate market, reducing overvaluation and undervaluation of properties. This contributes to a fairer and more transparent marketplace.

In conclusion, the application of machine learning in predicting house prices is a groundbreaking development with far-reaching implications. It empowers individuals, businesses, and governments to navigate the real estate market with more confidence and precision.